

# Chapter 3

## Transport Layer

A note on the use of these PowerPoint slides:


We're making these slides freely available to all (faculty, students, readers). They're in PowerPoint form so you see the animations; and can add, modify, and delete slides (including this one) and slide content to suit your needs. They obviously represent a *lot* of work on our part. In return for use, we only ask the following:

- If you use these slides (e.g., in a class) that you mention their source (after all, we'd like people to use our book!)
- If you post any slides on a www site, that you note that they are adapted from (or perhaps identical to) our slides, and note our copyright of this material.

For a revision history, see the slide note for this page.

Thanks and enjoy! JFK/KWR

All material copyright 1996-2023  
J.F Kurose and K.W. Ross, All Rights Reserved



## *Computer Networking: A Top-Down Approach*

8<sup>th</sup> edition

Jim Kurose, Keith Ross  
Pearson, 2020

# Transport layer: overview

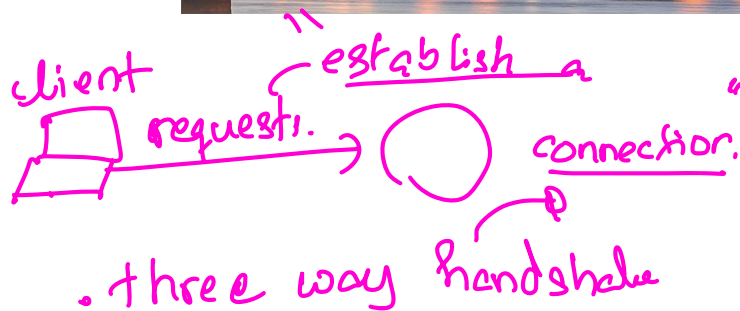
## *Our goal:*

- understand principles behind transport layer services:
  - multiplexing, demultiplexing
  - reliable data transfer
  - flow control
  - congestion control
- learn about Internet transport layer protocols:
  - UDP: connectionless transport
  - TCP: connection-oriented reliable transport
  - TCP congestion control

1M88 = 1460 Bytes if ethernet.

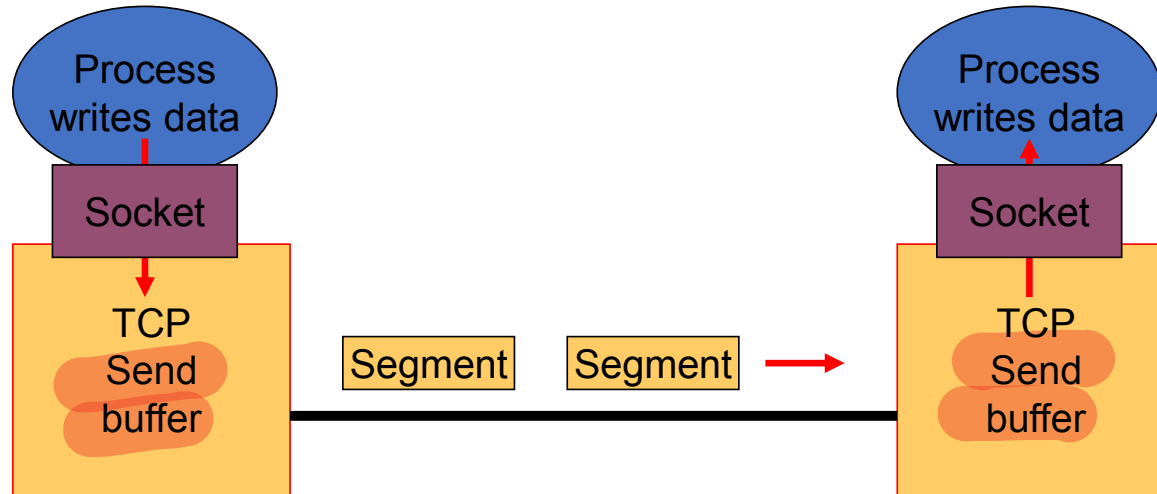
# Chapter 3: roadmap

- Transport-layer services
- Multiplexing and demultiplexing
- Connectionless transport: UDP
- Principles of reliable data transfer
- **Connection-oriented transport: TCP**
  - segment structure
  - reliable data transfer
  - flow control
  - connection management
- Principles of congestion control
- TCP congestion control



# TCP: overview

RFCs: 793, 1122, 2018, 5681, 7323



- The client process passes data through the socket.
- TCP directs data to the send's buffer. (memory allocate - pick up. storing station)
- TCP performs three-way handshake. (establish a connection)
- TCP sends data in segments.
- Segment sized is limited by the maximum segment size (MSS).

# TCP: overview

RFCs: 793, 1122, 2018, 5681, 7323

logical point to point connection.

## 1 point-to-point:

- one sender, one receiver

## 2 reliable, in-order byte stream:

- no "message boundaries"

## 3 full duplex data:

- bi-directional data flow in same connection
- MSS: maximum segment size

## 4 cumulative ACKs

## 5 pipelining:

- TCP congestion and flow control set window size

## 6 connection-oriented:

- handshaking (exchange of control messages) initializes sender, receiver state before data exchange

## 7 flow control:

- sender will not overwhelm receiver

3000 Bytes  
(6-2999)



1460 B  
0-1459

1460 B  
1460-2919



1460 B  
2920-2999

TCP  
↓

TCP Header 1460 Segment

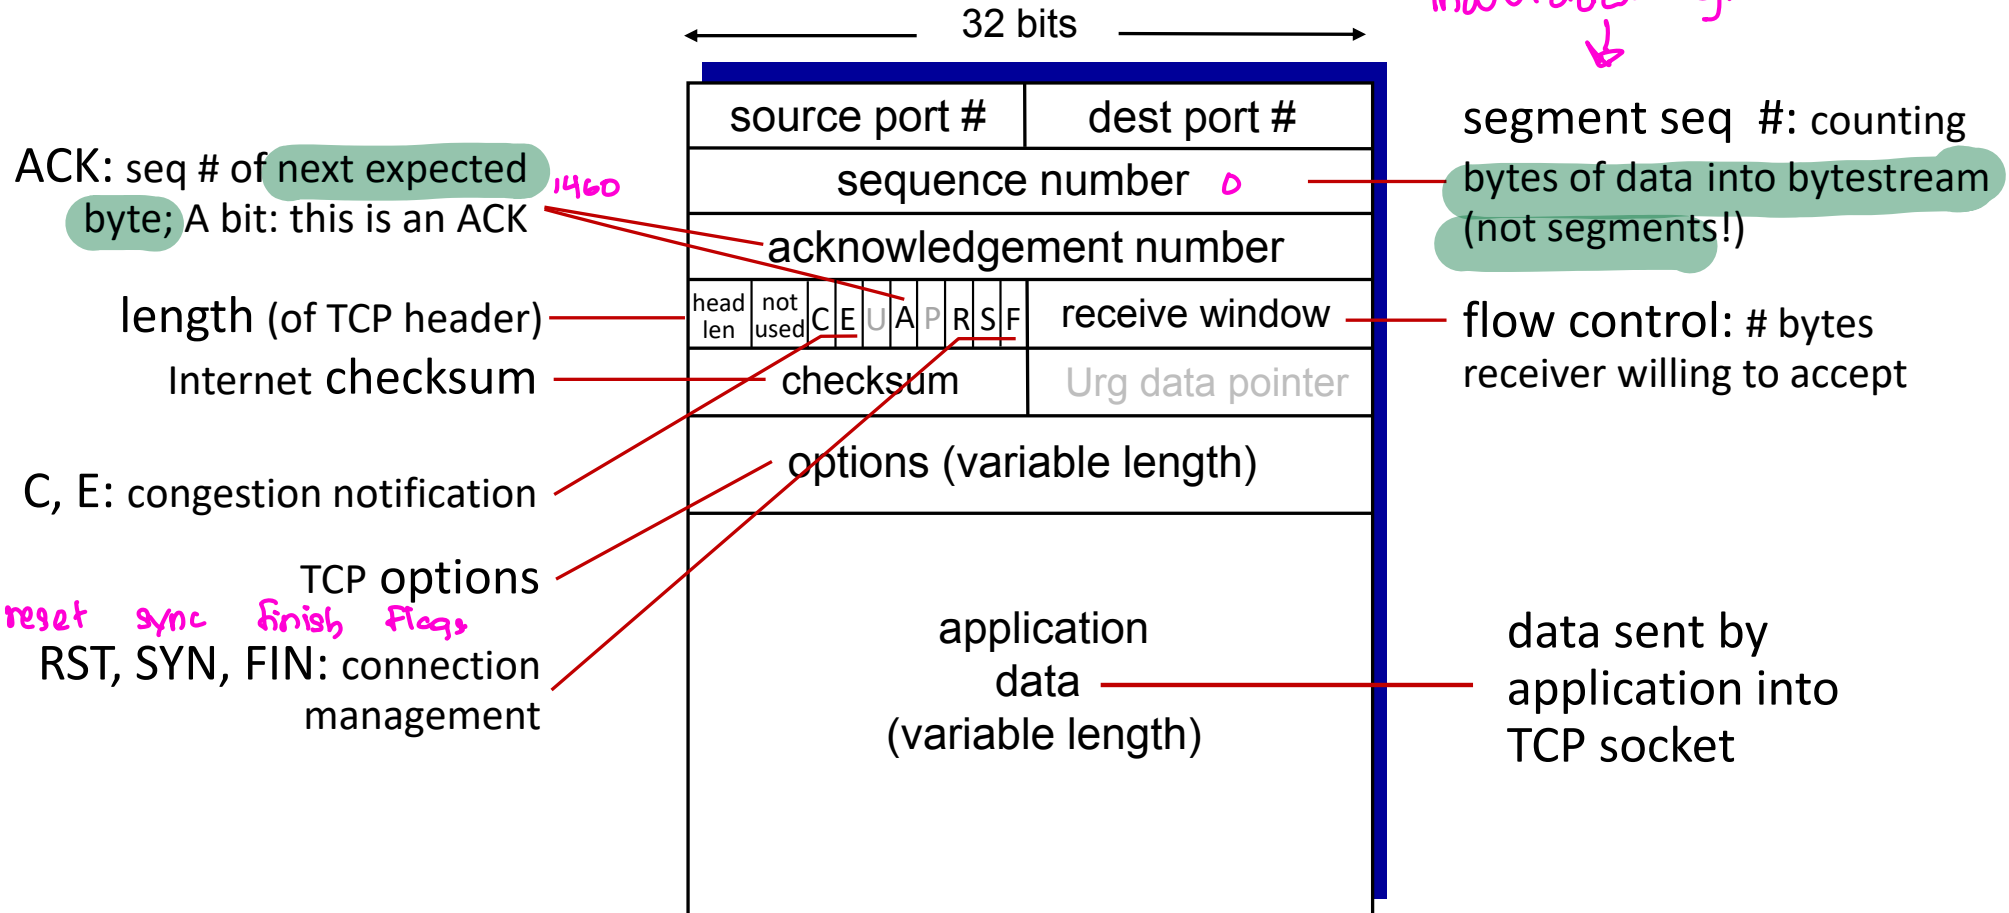
Full Duplex Data Bidirectional flow

Data #1



# TCP segment structure

0, 1460, 2920 seq numbers.  
first byte no of the individual bytestream  
↓



# TCP Sequence No. and Acknowledge No.

## Sequence numbers:


- TCP views data as unstructured, but ordered stream of bytes.
- Sequence numbers are over bytes, not segments
  - Byte stream number of first byte in segment's data
- Initial sequence number is chosen randomly
- TCP is full duplex – numbering of data is independent in each direction

## Acknowledgements:

- Acknowledgement number – sequence number of the next byte expected from the sender
- **ACKs are cumulative**

outgoing segment from sender

source port #	dest port #
sequence number	
acknowledgement number	
	rwnd
checksum	urg pointer



outgoing segment from receiver

source port #	dest port #
sequence number	
acknowledgement number	
	rwnd
checksum	urg pointer




# TCP Sequence No. and Acknowledge No.

Q: how receiver handles out-of-order segments

- A: TCP spec doesn't say, - up to implementor

outgoing segment from sender

source port #	dest port #
sequence number	
acknowledgement number	
	rwnd
checksum	urg pointer



outgoing segment from receiver

source port #	dest port #
sequence number	
acknowledgement number	
	rwnd
checksum	urg pointer

# TCP sequence numbers, ACKs

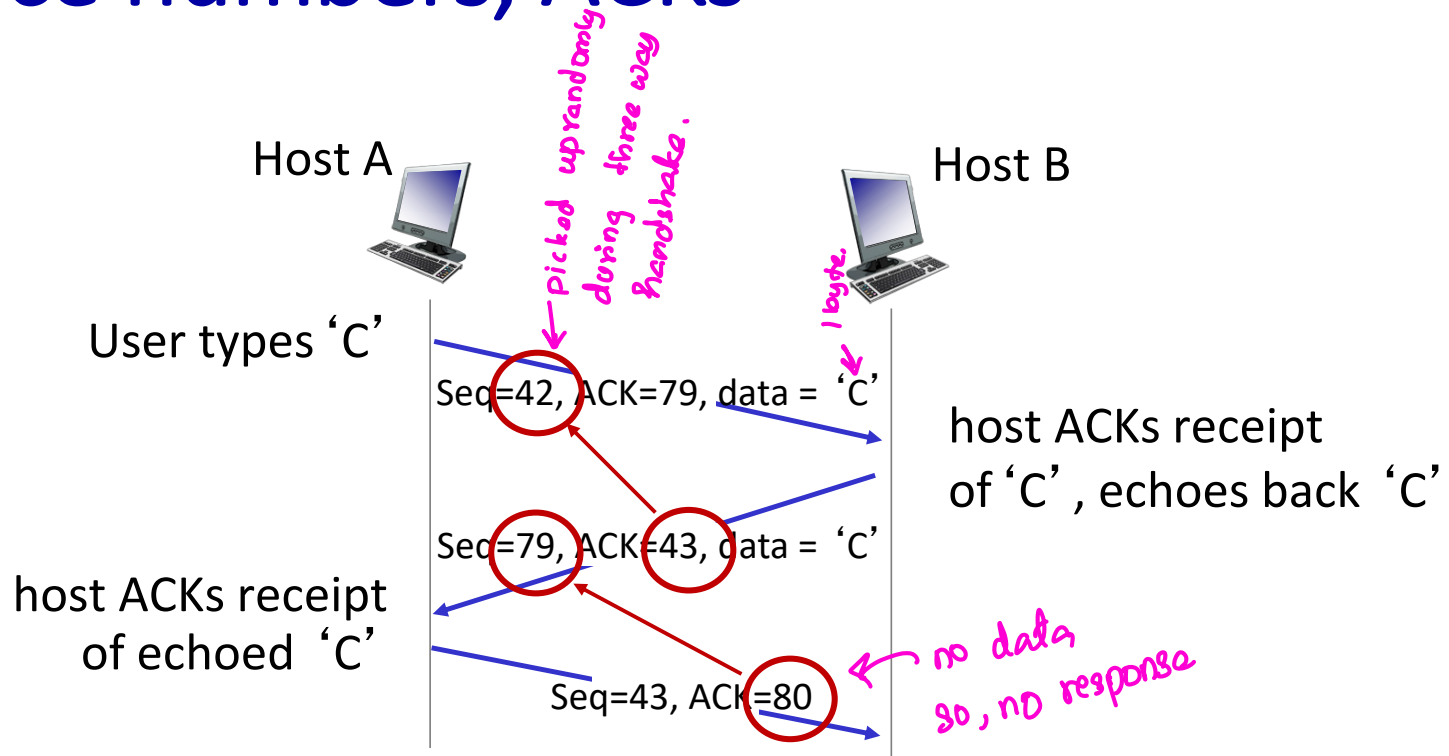
## Seq. numbers:

- byte stream  
“number” of first  
byte in segment’s  
data

## ACKs:

- seq # of next byte  
expected from  
other side
- **cumulative ACK**

**Q: how receiver  
handles out-of-  
order ?**




simple telnet scenario

# TCP round trip time, timeout

round trip time may vary.

Q: how to set TCP timeout value?

- longer than RTT, but RTT varies!
- *too short*: premature timeout, unnecessary retransmissions
- *too long*: slow reaction to segment loss




Q: how to estimate RTT?

- `SampleRTT`: measured time from segment transmission until ACK receipt
  - ignore retransmissions
- `SampleRTT` will vary, want estimated RTT “smoother”
  - average several *recent* measurements, not just current `SampleRTT`

# ? TCP round trip time, timeout

$$\text{EstimatedRTT} = (1 - \alpha) * \text{EstimatedRTT} + \alpha * \text{SampleRTT}$$

- exponential weighted moving average (EWMA)
- influence of past sample decreases exponentially fast
- typical value:  $\alpha = 0.125$



# ? TCP round trip time, timeout

- timeout interval: **EstimatedRTT** plus “safety margin”
  - large variation in **EstimatedRTT**: want a larger safety margin

$$\text{TimeoutInterval} = \text{EstimatedRTT} + 4 * \text{DevRTT}$$

 estimated RTT      “safety margin”

*(est rt - current?)*


- **DevRTT**: EWMA of **SampleRTT** deviation from **EstimatedRTT**:

$$\text{DevRTT} = (1 - \beta) * \text{DevRTT} + \beta * |\text{SampleRTT} - \text{EstimatedRTT}|$$

(typically,  $\beta = 0.25$ )

# TCP Reliable Data Transfer

- TCP creates **rdt** service on top of IP's unreliable service
  - Pipelined segments (multiple UnACK'd in pipes)
  - Cumulative ACKs
  - Single retransmission timer
- Retransmissions are triggered by:
  - timeout events
  - duplicate ACKs



Initially consider simplified TCP sender:

- ignore duplicate ACKs
- ignore flow control, congestion control

# TCP Sender Events (1)

## data rcvd from app:

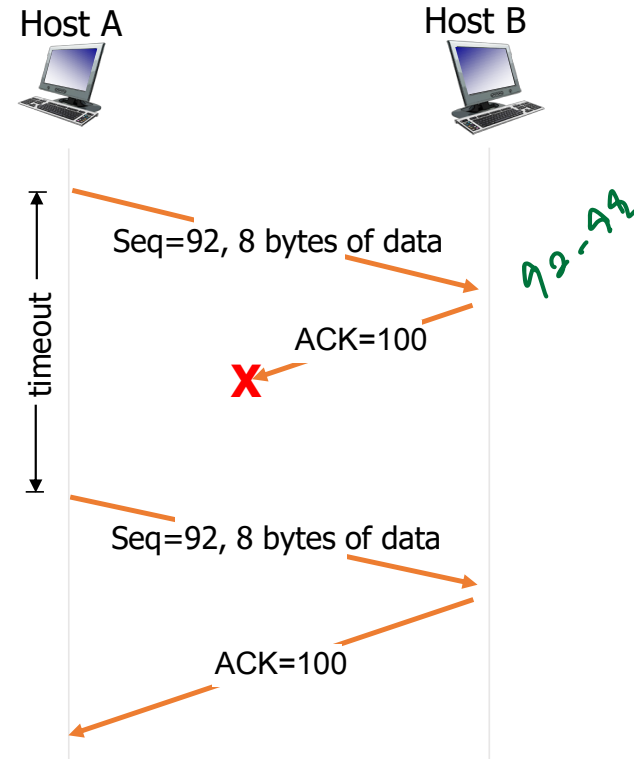
- Create segment with seq #
- seq # is byte-stream number of first data byte in segment
- **start timer** if not already running (think of timer as for oldest unacked segment) *the oldest one triggered*
- expiration interval: `TimeoutInterval`

## timeout:

- retransmit segment that caused timeout
- restart timer

## ack rcvd:

- If acknowledges previously unacked segments
  - update what is known to be acked
  - start timer if there are outstanding segments



Retransmission due to a lost acknowledgement

# TCP Sender Events (2)

## data rcvd from app:

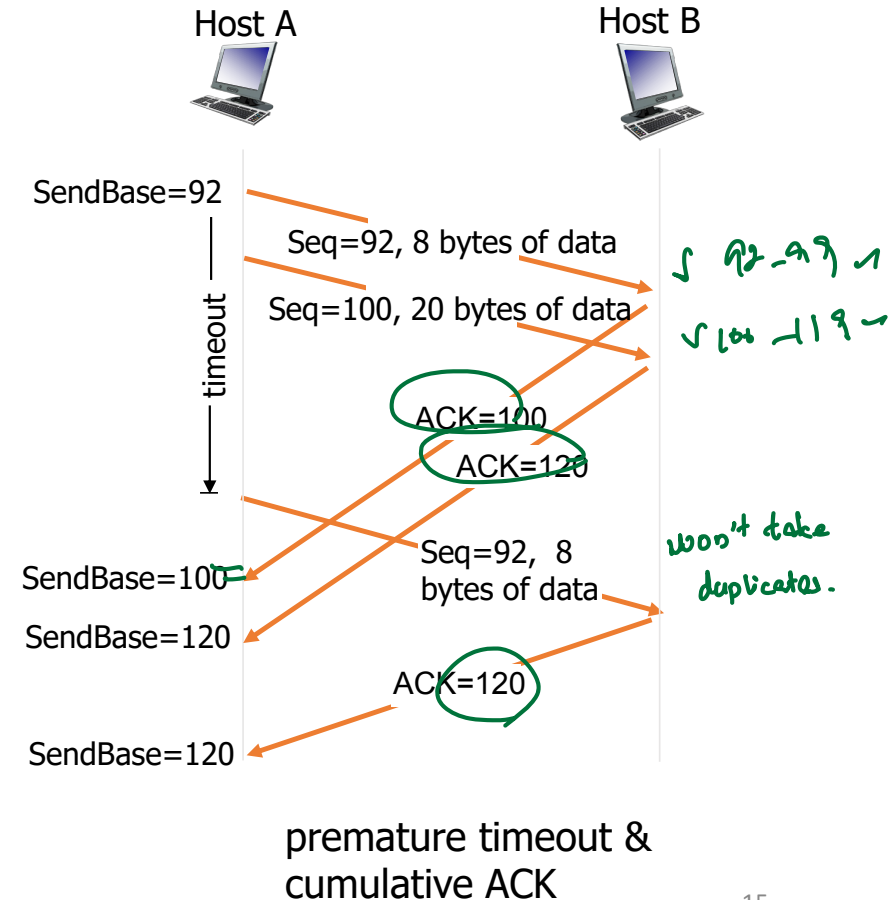
- Create segment with seq #
- seq # is byte-stream number of first data byte in segment
- **start timer** if not already running (think of timer as for oldest unacked segment)
- expiration interval: `TimeoutInterval`

## timeout:

- retransmit segment that caused timeout
- restart timer

## ack rcvd:

- If acknowledges previously unacked segments
  - update what is known to be acked
  - start timer if there are outstanding segments





# TCP Sender Events (3)

## data rcvd from app:

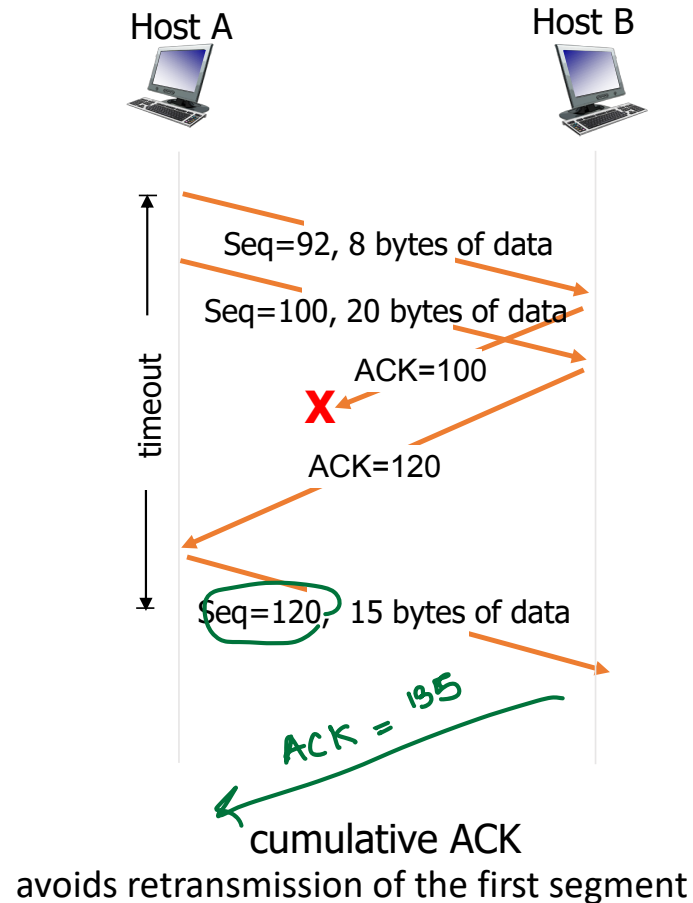
- Create segment with seq #
- seq # **is** byte-stream number of first data byte in segment
- **start timer** if not already running (think of timer as for oldest unacked segment)
- expiration interval: `TimeoutInterval`

## timeout:



- retransmit segment that caused timeout
- restart timer

## ack rcvd:

- If acknowledges previously unacked segments
  - update what is known to be acked
  - start timer if there are outstanding segments



# TCP ACK Generation [RFC 1122, 2581, 5681, 7323]



## event at receiver

## TCP receiver action

arrival of in-order segment with expected seq #. All data up to expected seq # already ACKed

delayed ACK. Wait up to 500ms for next segment. If no next segment, send ACK

arrival of in-order segment with expected seq #. One other segment has ACK pending

immediately send single cumulative ACK, ACKing both in-order segments

arrival of out-of-order segment higher-than-expect seq. # . Gap detected

immediately send **duplicate ACK**, indicating seq. # of next expected byte

arrival of segment that partially or completely fills gap

immediate send ACK, provided that segment starts at lower end of gap

# TCP fast retransmit

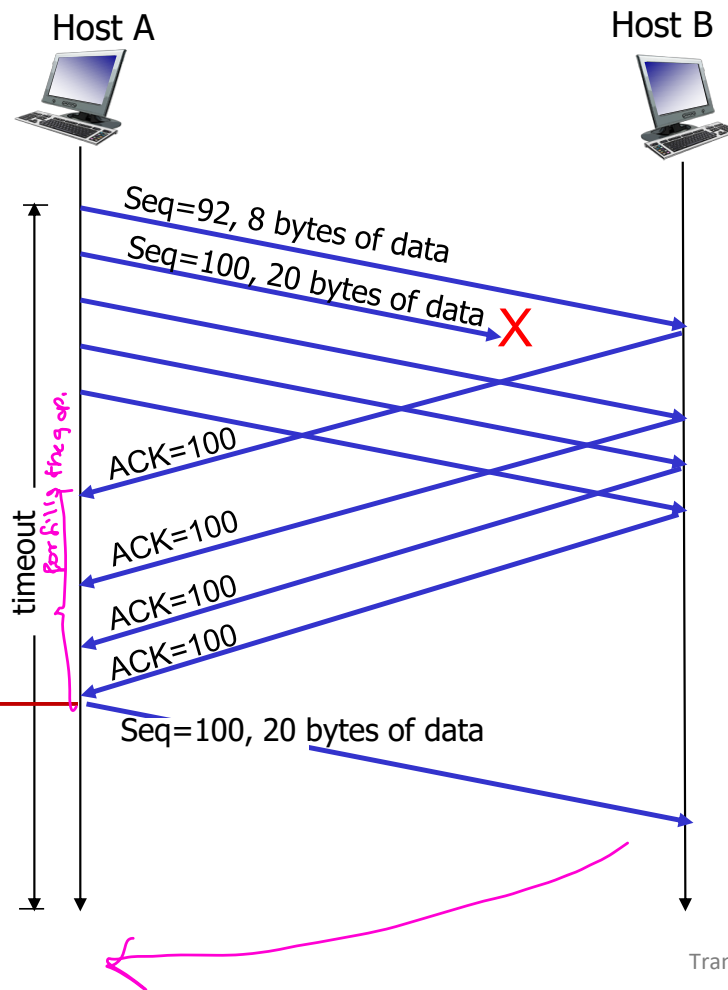
## *TCP fast retransmit*

if sender receives **3 additional ACKs** for same data (“triple duplicate ACKs”), resend unACKed segment with smallest seq #

- likely that unACKed segment lost, so don't wait for timeout

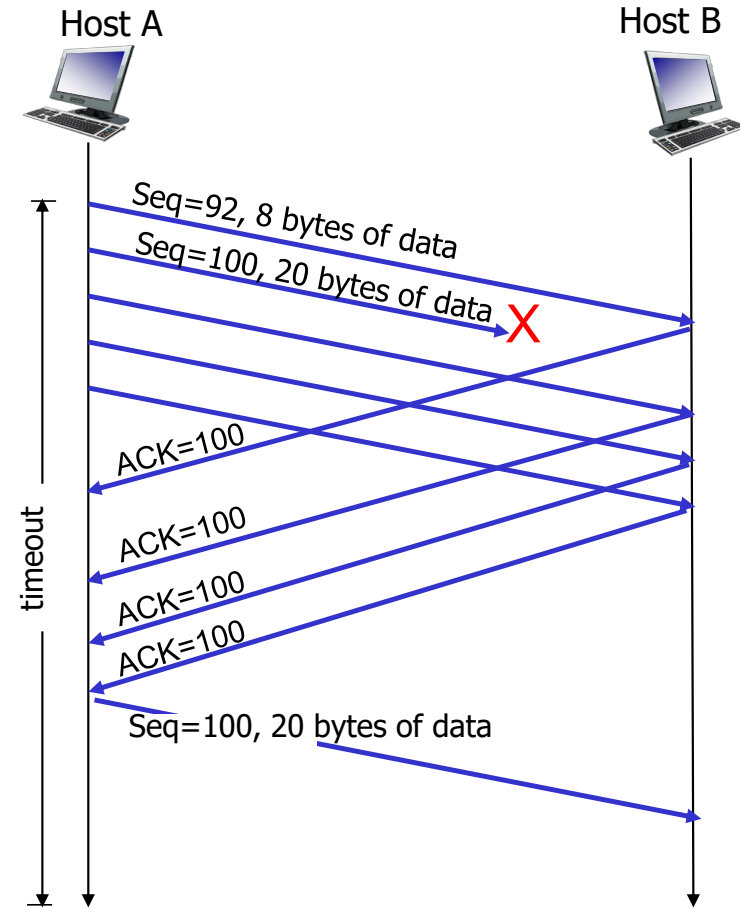


Receipt of three duplicate ACKs indicates 3 segments received after a missing segment – lost segment is likely. So retransmit!



# TCP fast retransmit

- Since the timeout interval is exponentially increased, increasing end-to-end delay.
  - Long delay before retransmission\*
- Fortunately, the sender can often detect packet loss well before the timer expires by just **duplicate ACKs**.
  - sender often sends many segments back-to-back (send many segments one after another)
  - if one segment is lost, there will likely be many **duplicate ACKs**.
- If the TCP sender receives **three duplicate ACKs** for the same data, TCP performs a **fast retransmit**, retransmitting the missing segment before the timer expires.



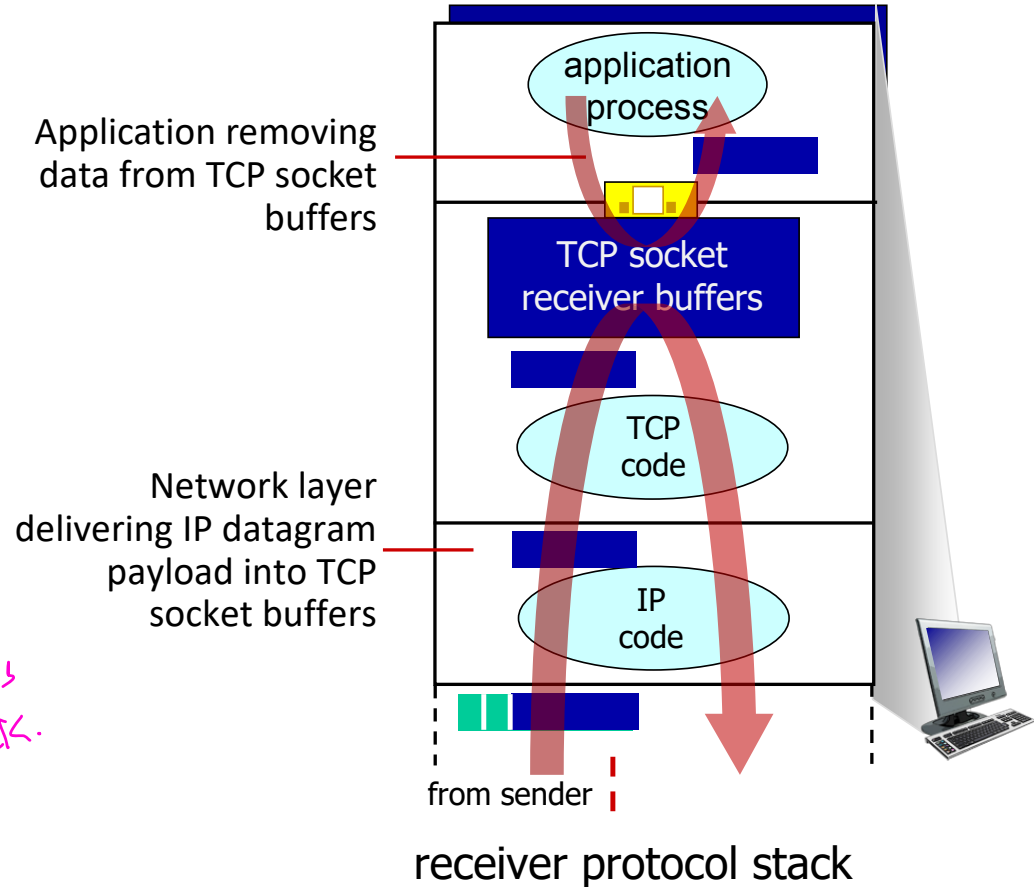

# Chapter 3: roadmap

- Transport-layer services
- Multiplexing and demultiplexing
- Connectionless transport: UDP
- Principles of reliable data transfer
- **Connection-oriented transport: TCP**
  - segment structure
  - reliable data transfer
  - flow control
  - connection management
- Principles of congestion control
- TCP congestion control



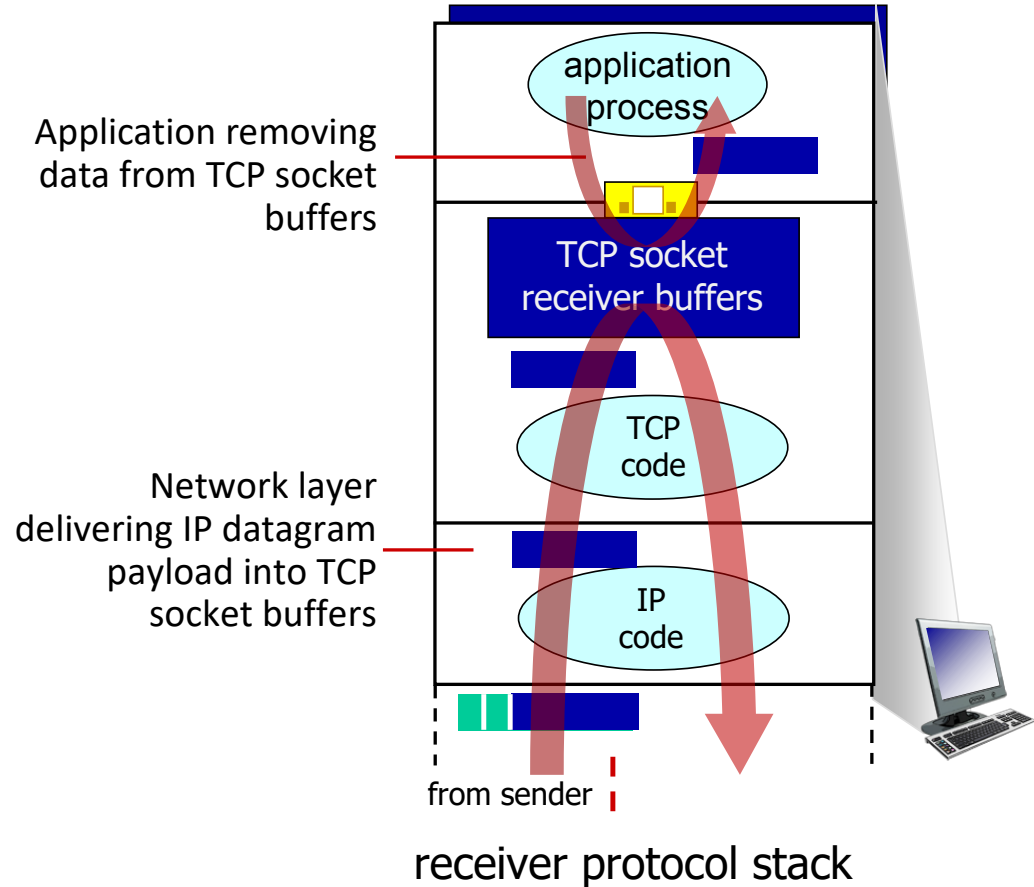
# TCP flow control

Q: What happens if network layer delivers data faster than application layer removes data from socket buffers (receive buffers)?



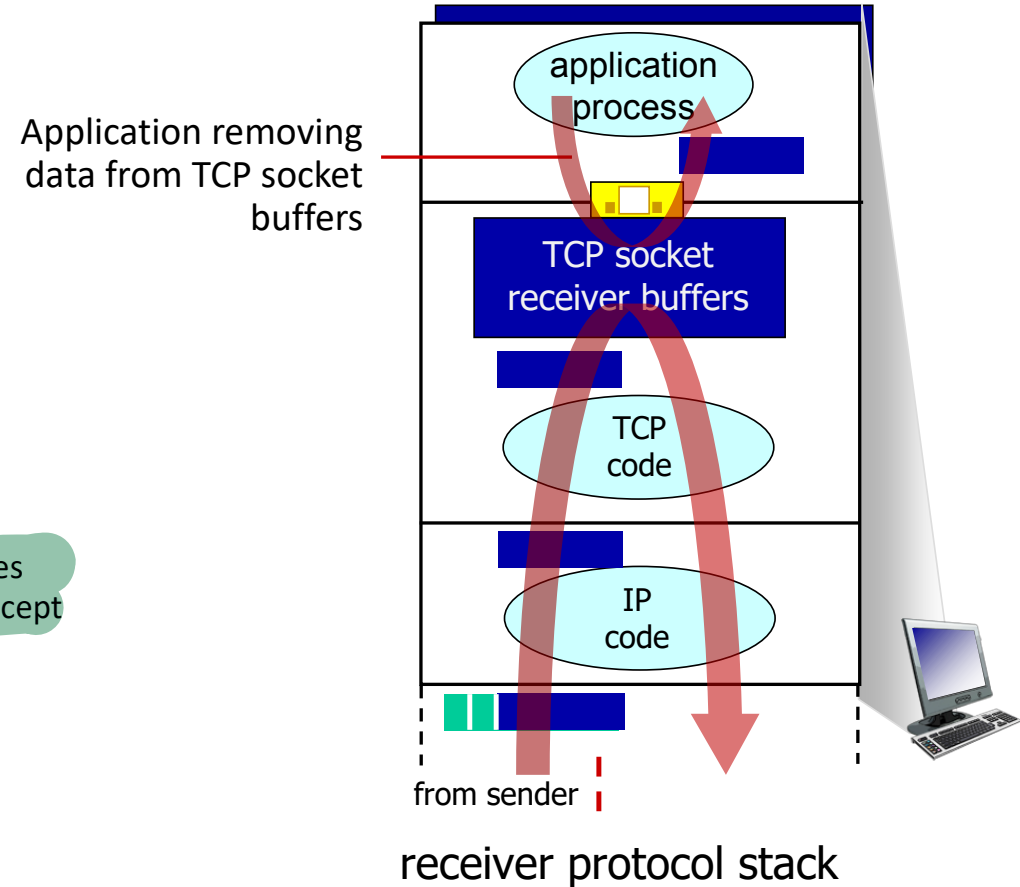
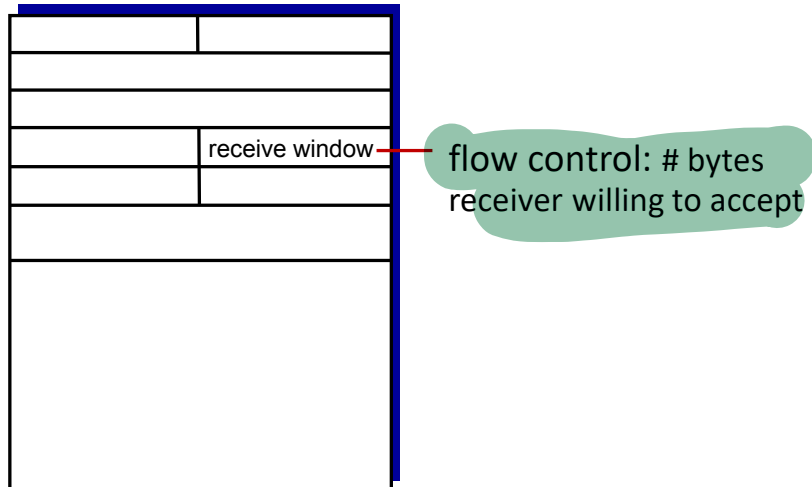
# TCP flow control

Q: What happens if network layer delivers data faster than application layer removes data from socket buffers (receive buffers)?



# TCP flow control

Q: What happens if network layer delivers data faster than application layer removes data from socket buffers (receive buffers)?



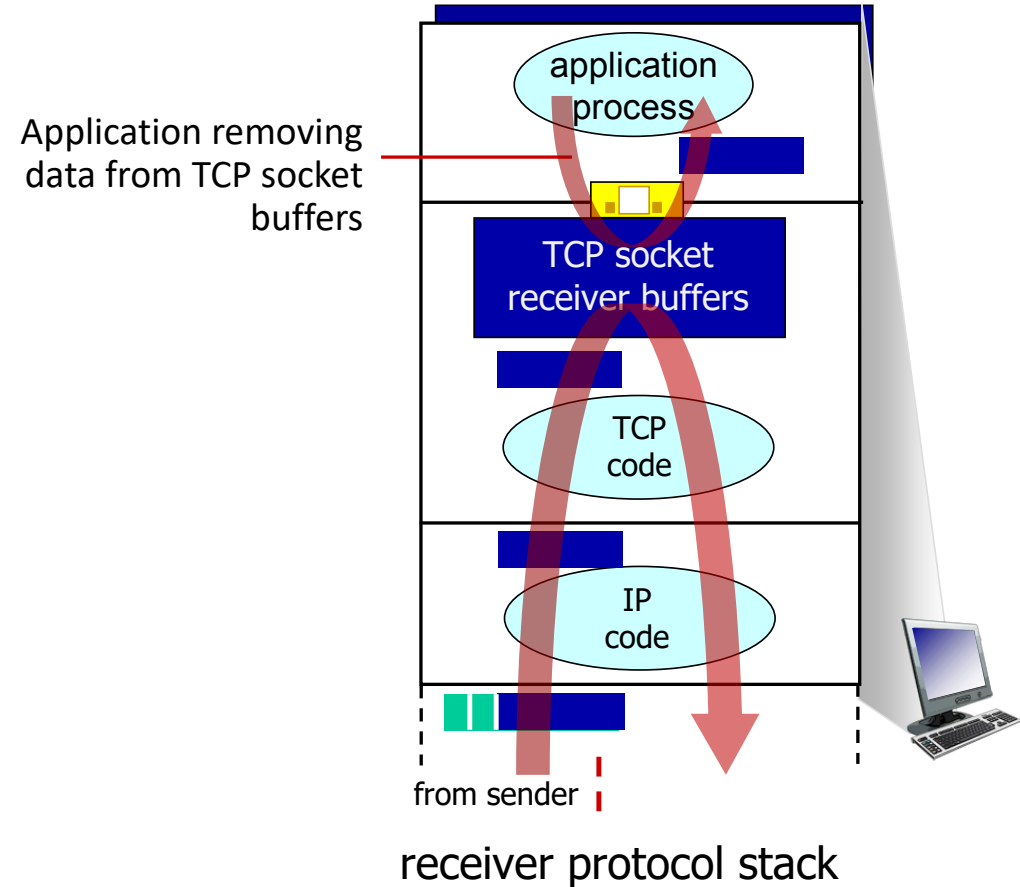


# TCP flow control

Q: What happens if network layer delivers data faster than application layer removes data from socket buffers (receive buffers)?

## flow control


receiver controls sender, so sender won't overflow receiver's buffer by transmitting too much, too fast



# TCP flow control

clap.

- TCP receiver “advertises” **free buffer space** in **rwnd** field in TCP header
  - **RcvBuffer** size set via socket options (typical default is 4096 bytes)
  - many operating systems auto-adjust **RcvBuffer**
- sender limits amount of unACKed (“in-flight”) data to received **rwnd**
- guarantees receive buffer will not overflow



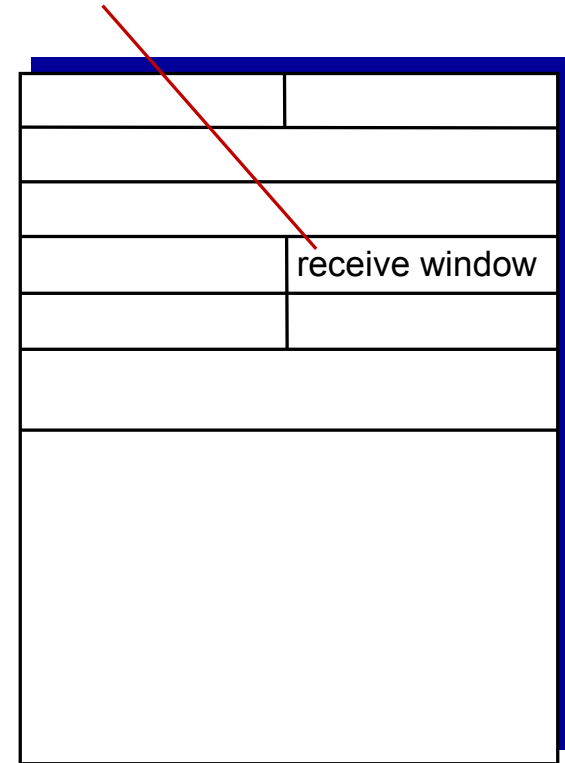
TCP receiver-side buffering

# TCP flow control

Shuf

- TCP receiver “advertises” **free buffer space** in **rwnd** field in TCP header
  - **RcvBuffer** size set via socket options (typical default is 4096 bytes)
  - many operating systems auto-adjust **RcvBuffer**
- sender limits amount of unACKed (“in-flight”) data to received **rwnd**
- guarantees receive buffer will not overflow

flow control: # bytes receiver willing to accept




TCP segment format


# TCP connection management <sup>sketch</sup>

before exchanging data, sender/receiver “handshake”:

- agree to establish connection (each knowing the other willing to establish connection)
- agree on connection parameters (**starting seq #** and **rwnd**)



```
Socket clientSocket =  
    newSocket("hostname", "port number");
```



```
Socket connectionSocket =  
    welcomeSocket.accept();
```

(review)

# TCP Connection Management

## 3-Way Handshake:


Step 1: client host sends TCP SYN segment to server

- specifies initial seq #
- no data

Step 2: server host receives SYN, (if want to communicate) replies with SYN/ACK segment

- server allocates buffers
- specifies server initial seq. #

Step 3: client receives SYN/ACK, replies with ACK segment, which may contain data



# TCP 3-Way Handshake

*client state*

LISTEN

SYNSENT

ESTAB

choose init seq num, x  
send TCP SYN msg

SYNbit=1, Seq=x

SYNbit=1, Seq=y  
ACKbit=1; ACKnum=x+1

received SYNACK(x)  
indicates server is live;  
send ACK for SYNACK;  
this segment may contain  
client-to-server data

ACKbit=1, ACKnum=y+1



*server state*

LISTEN

SYN RCVD

ESTAB

choose init seq num, y  
send TCP SYNACK  
msg, acking SYN

received ACK(y)  
indicates client is live

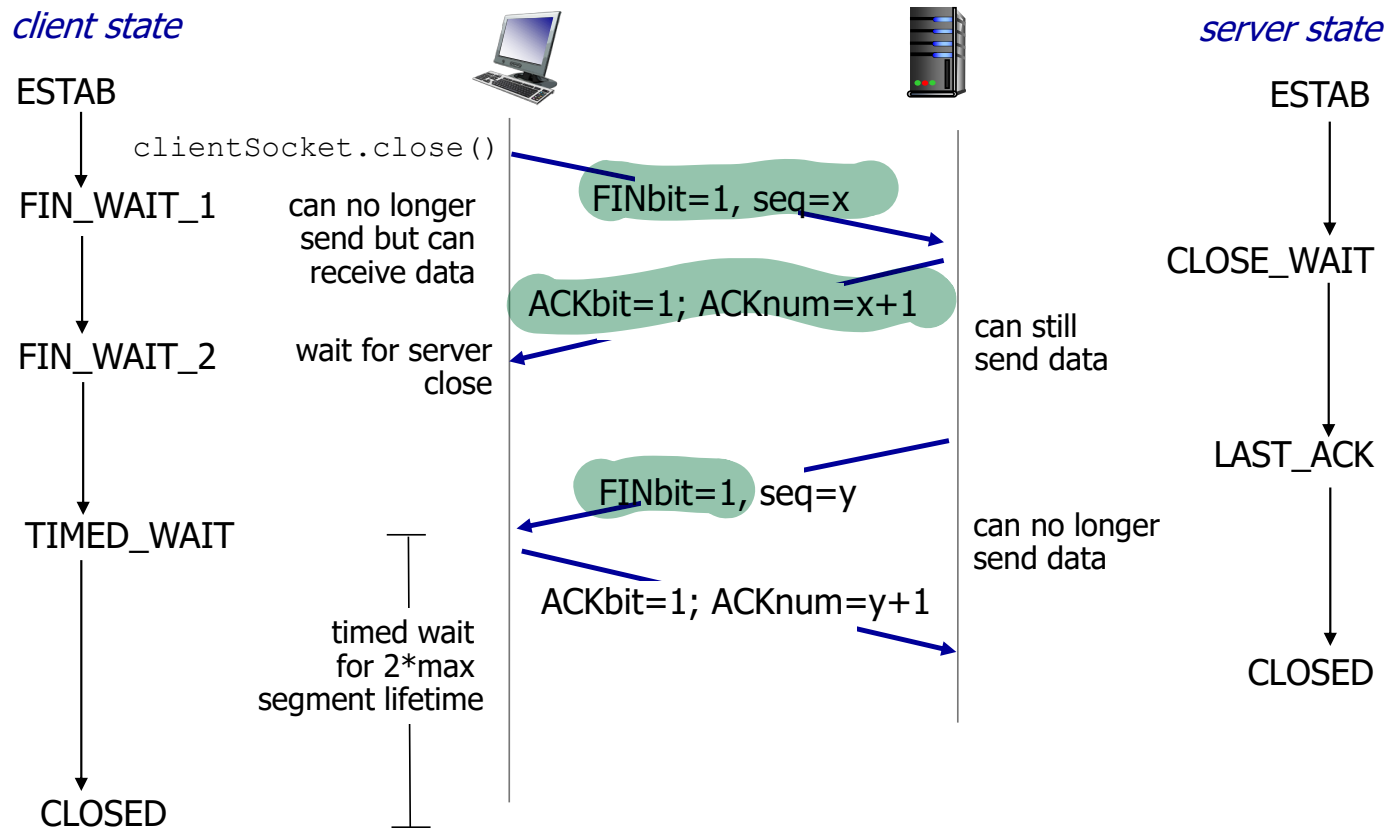
# A human 3-way handshake protocol



(review!)

# TCP Connection Termination

- client, server each close their side of connection
- send TCP segment with FIN bit = 1
- respond to received FIN with ACK
- on receiving FIN, ACK can be combined with own FIN
- simultaneous FIN exchanges can be handled





# Chapter 3: roadmap

- Transport-layer services
- Multiplexing and demultiplexing
- Connectionless transport: UDP
- Principles of reliable data transfer
- Connection-oriented transport: TCP
- **Principles of congestion control**
- TCP congestion control
- Evolution of transport-layer functionality



# Principles of congestion control

(slows down sending rate of the sender?)  
observes network condition.

## Congestion:

- informally: “too many sources sending too much data too fast for **network** to handle”
- manifestations:
  - long delays (queueing in router buffers)
  - packet loss (buffer overflow at routers)
- different from flow control!
- a top-10 problem!

buffer full, bandwidth full



### congestion control:

too many senders,  
sending too fast

**flow control:** one sender  
too fast for one receiver



one sender  
and not one flow  
in receiver

# TCP Congestion Control: Overview

- TCP uses end-to-end congestion control.
- It limits the sender's sending rate.
- If the sender perceives that there is **little (no) congestion** on the path, the TCP sender **increases its send rate**.
- If the sender perceives that there is **congestion** on the path, the TCP sender **reduces its send rate**.





main actors → router



# Congestion Control: Approaches

- **Goal:** Throttle senders as needed to ensure load on the network is “reasonable”
- **End-end congestion control:**
  - no explicit feedback from network
  - congestion inferred from end-system observed loss, delay
  - approach taken by TCP
- **Network-assisted congestion control:**
  - routers provide feedback to end systems
  - single bit indicating congestion
  - explicit rate sender should send at
  - TCP ECN, ATM, DECbit protocols



# Chapter 3: roadmap

- Transport-layer services
- Multiplexing and demultiplexing
- Connectionless transport: UDP
- Principles of reliable data transfer
- Connection-oriented transport: TCP
- Principles of congestion control
- **TCP congestion control**
- **Evolution of transport-layer functionality**



# TCP congestion control: AIMD

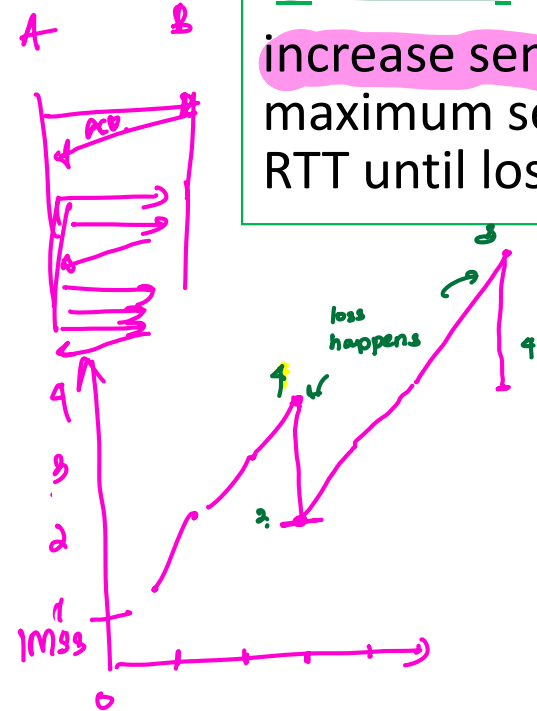
- *approach*: senders can increase sending rate until packet loss (congestion) occurs, then decrease sending rate on loss event

## Additive Increase


increase sending rate by 1 maximum segment size every RTT until loss detected

## Multiplicative Decrease

cut sending rate in half at each loss event



TCP sender Sending rate



**AIMD** sawtooth behavior: *probing* for bandwidth

# TCP AIMD: more


*Multiplicative decrease* detail: sending rate is

- Cut in half on loss detected by triple duplicate ~~ACK (TCP Reno)~~
- Cut to 1 MSS (maximum segment size) when loss detected by timeout (~~TCP Tahoe~~)

Why AIMD?

- AIMD – a distributed, asynchronous algorithm – has been shown to:
  - optimize congested flow rates network wide!
  - have desirable stability properties

# TCP congestion control: details



## TCP sending behavior:

- *roughly*: send `cwnd` bytes, wait RTT for ACKS, then send more bytes


$$\text{TCP rate} \approx \frac{\text{cwnd}}{\text{RTT}} \text{ bytes/sec}$$

- TCP sender limits transmission:  $\text{LastByteSent} - \text{LastByteAcked} \leq \text{cwnd}$
- `cwnd` is dynamically adjusted in response to observed network congestion (implementing TCP congestion control)



# TCP slow start

- when connection begins, increase rate **exponentially** until first loss event:
  - initially **cwnd = 1 MSS**
  - **double cwnd** every RTT
  - done by incrementing **cwnd** for every ACK received
- **summary:** initial rate is slow, but ramps up exponentially fast




# TCP: from Slow Start to Congestion Avoidance

**Q:** when should the exponential increase switch to linear?

**A:** when **cwnd** gets to 1/2 of its value before timeout.

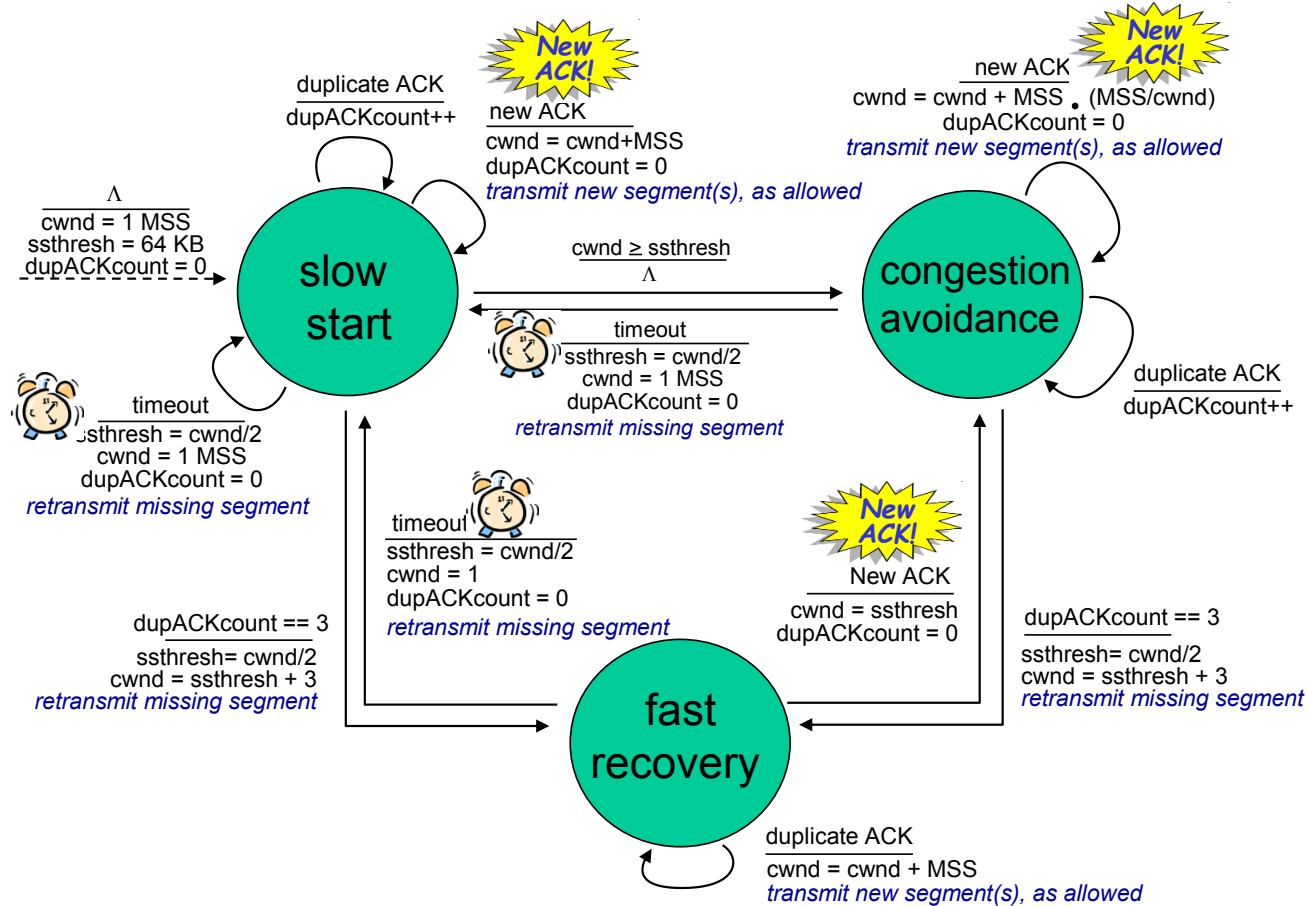
## Implementation:

- variable **ssthresh**
- on loss event, **ssthresh** is set to 1/2 of **cwnd** just before loss event



\* Check out the online interactive exercises for more examples: [http://gaia.cs.umass.edu/kurose\\_ross/interactive/](http://gaia.cs.umass.edu/kurose_ross/interactive/)

# Summary: TCP congestion control




# TCP CUBIC

- Is there a better way than AIMD to “probe” for usable bandwidth?


- Insight/intuition:

- $W_{\max}$ : sending rate at which congestion loss was detected
- congestion state of bottleneck link probably (?) hasn't changed much
- after cutting rate/window in half on loss, initially ramp to to  $W_{\max}$  *faster*, but then approach  $W_{\max}$  more *slowly*




# TCP CUBIC

- K: point in time when TCP window size will reach  $W_{\max}$ 
  - K itself is tunable
- increase  $W$  as a function of the *cube* of the distance between current time and K
  - larger increases when further away from K
  - smaller increases (cautious) when nearer K
- TCP CUBIC default in Linux, most popular TCP for popular Web servers



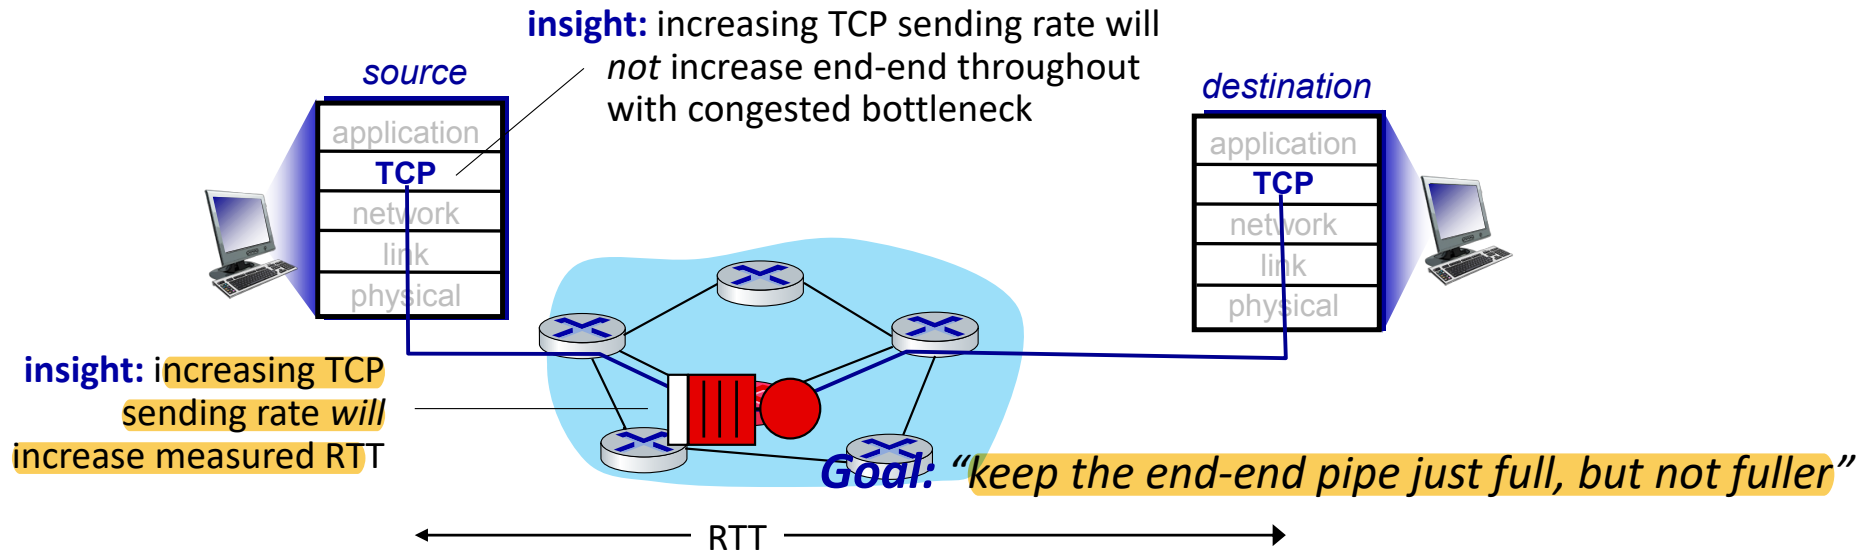
# TCP and the congested “bottleneck link”

- TCP (classic, CUBIC) increase TCP's sending rate until packet loss occurs at some router's output: the *bottleneck link*



# TCP and the congested “bottleneck link”


- TCP (classic, CUBIC) increase TCP’s sending rate until packet loss occurs at some router’s output: the *bottleneck link*
- understanding congestion: useful to focus on congested bottleneck link



Don't want until loss, Act when delay happens.

# Delay-based TCP congestion control

Keeping sender-to-receiver pipe “just full enough, but no fuller”: keep bottleneck link busy transmitting, but avoid high delays/buffering



$$\text{measured throughput} = \frac{\text{\# bytes sent in last RTT interval}}{\text{RTT}_{\text{measured}}}$$

*round trip*

## Delay-based approach:

- $\text{RTT}_{\min}$  - minimum observed RTT (uncongested path)
- uncongested throughput with congestion window  $\text{cwnd}$  is  $\text{cwnd}/\text{RTT}_{\min}$ 
  - ✓ if measured throughput “very close” to uncongested throughput  
increase  $\text{cwnd}$  linearly /\* since path not congested \*/
  - ✓ else if measured throughput “far below” uncongested throughput  
decrease  $\text{cwnd}$  linearly /\* since path is congested \*/




# Delay-based TCP congestion control

- **congestion control without inducing/forcing loss**
- maximizing throughput (“keeping the just pipe full...”) while keeping delay low (“...but not fuller”)
- a number of deployed TCPs take a delay-based approach
  - BBR deployed on Google’s (internal) backbone network

# Explicit congestion notification (ECN)

TCP deployments often implement *network-assisted* congestion control:


- two bits in IP header (ToS field) marked *by network router* to indicate congestion
  - *policy* to determine marking chosen by network operator
- congestion indication carried to destination
- destination sets ECE bit on ACK segment to notify sender of congestion
- involves both IP (IP header ECN bit marking) and TCP (TCP header C,E bit marking)



# TCP fairness

no fairness between different TCP connections


**Fairness goal:** if  $K$  TCP sessions share same bottleneck link of bandwidth  $R$ , each should have average rate of  $R/K$



# Q: is TCP Fair?

Example: two competing TCP sessions:

- additive increase gives slope of 1, as throughput increases
- multiplicative decrease decreases throughput proportionally



*Is TCP fair?*

**A:** Yes, under idealized assumptions:

- same RTT
- fixed number of sessions only in congestion avoidance


# Fairness: must all network apps be “fair”?

## Fairness and UDP

- multimedia apps often do not use TCP
  - do not want rate throttled by congestion control
- instead use UDP:
  - send audio/video at constant rate, tolerate packet loss
- there is no “Internet police” policing use of congestion control

## Fairness, parallel TCP connections

- application can open *multiple* parallel connections between two hosts
- web browsers do this , e.g., link of rate  $R$  with 9 existing connections:
  - new app asks for 1 TCP, gets rate  $R/10$
  - new app asks for 11 TCPs, gets  $R/2$



→ UDP has more benefit over net work.

# Transport layer: roadmap

- Transport-layer services
- Multiplexing and demultiplexing
- Connectionless transport: UDP
- Principles of reliable data transfer
- Connection-oriented transport: TCP
- Principles of congestion control
- TCP congestion control
- Evolution of transport-layer functionality



# Evolving transport-layer functionality *def.*

- TCP, UDP: principal transport protocols for 40 years
- different “flavors” of TCP developed, for specific scenarios:

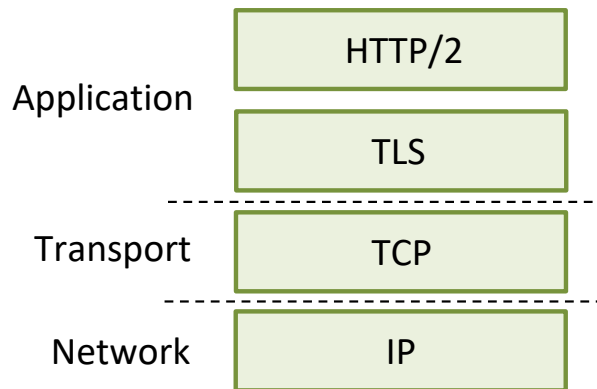
Scenario	Challenges
Long, fat pipes (large data transfers)	Many packets “in flight”; loss shuts down pipeline
Wireless networks	Loss due to noisy wireless links, mobility; TCP treat this as congestion loss
Long-delay links	Extremely long RTTs
Data center networks	Latency sensitive
Background traffic flows	Low priority, “background” TCP flows

- moving transport-layer functions to application layer, on top of UDP
  - HTTP/2 & HTTP/3: QUIC

research from

# QUIC: Quick UDP Internet Connections

- application-layer protocol, on top of UDP
  - increase performance of HTTP
  - deployed on many Google servers, apps (Chrome, mobile YouTube app)



HTTP/2 over TCP



# QUIC: Quick UDP Internet Connections

adopts approaches we've studied in this chapter for connection establishment, error control, congestion control

- **error and congestion control:** “Readers familiar with TCP’s loss detection and congestion control will find algorithms here that parallel well-known TCP ones.” [from QUIC specification]
- **connection establishment:** reliability, congestion control, authentication, encryption, state established in one RTT
- multiple application-level “streams” multiplexed over single QUIC connection
  - separate reliable data transfer, security
  - common congestion control

# Chapter 3: summary

- principles behind transport layer services:
  - multiplexing, demultiplexing
  - reliable data transfer
  - flow control
  - congestion control
- instantiation, implementation in the Internet
  - UDP
  - TCP

## Up next:

- leaving the network “edge” (application, transport layers)
- into the network “core”
- two network-layer chapters:
  - data plane
  - control plane